

Project-1

December 20, 2023

1 Course End Project : HealthCare

```
[1]: # importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

2 1. Data preprocessing and Visualisation

```
[2]: df = pd.read_csv("1645792390_cep1_dataset.csv")
```

```
[3]: df.head()
```

```
[3]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[4]: df.tail()
```

```
[4]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	
301	57	1	0	130	131	0	1	115	1	1.2	
302	57	0	1	130	236	0	0	174	0	0.0	

	slope	ca	thal	target
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

```
[5]: df.dtypes
```

```
[5]: age          int64
sex           int64
cp           int64
trestbps     int64
chol         int64
fbs          int64
restecg      int64
thalach      int64
exang        int64
oldpeak      float64
slope        int64
ca           int64
thal         int64
target       int64
dtype: object
```

```
[6]: df.shape
```

```
[6]: (303, 14)
```

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   int64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   int64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
```

```

10  slope      303 non-null    int64
11  ca         303 non-null    int64
12  thal       303 non-null    int64
13  target     303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

```

```
[8]: df.isnull().sum()    # there is no null values
```

```

[8]: age          0
     sex          0
     cp          0
     trestbps     0
     chol         0
     fbs          0
     restecg      0
     thalach      0
     exang        0
     oldpeak      0
     slope        0
     ca           0
     thal         0
     target       0
     dtype: int64

```

```
[9]: df[df.duplicated()]    # row number 164 is a duplicate row
```

```

[9]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
164    38    1   2         138   175    0         1       173     0       0.0

      slope  ca  thal  target
164      2   4     2        1

```

```

[10]: df.drop_duplicates(inplace=True)
      df.reset_index(drop=True, inplace=True)    # removing the duplicate row
      df.shape

```


```
[10]: (302, 14)
```

```

[11]: # checking for outliers
      plt.figure(figsize=(10,6))
      df.boxplot(column=['age', 'trestbps', 'chol', 'thalach', 'oldpeak'])

```

```
[11]: <Axes: >
```



```
[12]: # there are outliers in the numerical columns except "age"
```

```
[13]: # Finding values of outliers(IQR method)
```

```
def detect_outliers_iqr(data):
    outlier_list = []
    data = sorted(data)
    q1 = np.percentile(data, 25)
    q3 = np.percentile(data, 75)
    #print("The Val of Q1 and Q2",q1, q3)
    IQR = q3-q1
    lwr_bound = q1-(1.5*IQR)
    upr_bound = q3+(1.5*IQR)
    #print("The lower & Upper Bound",lwr_bound, upr_bound)

    for i in data:
        if (i<lwr_bound or i>upr_bound):
            outlier_list.append(i)
    return outlier_list # Driver code

for i in ['age','trestbps','chol','thalach','oldpeak']:
    outliers = detect_outliers_iqr(df[i])
    print("Outliers in",i,"attribute :", outliers)
```

```

Outliers in age attribute : []
Outliers in trestbps attribute : [172, 174, 178, 178, 180, 180, 180, 192, 200]
Outliers in chol attribute : [394, 407, 409, 417, 564]
Outliers in thalach attribute : [71]
Outliers in oldpeak attribute : [4.2, 4.2, 4.4, 5.6, 6.2]

```

```

[14]: # Handling outliers using Quantile based flooring and capping method.
#the outlier is capped at a certain value above the 90th percentile value or
      ↪floored at a factor below the 10th percentile value
# Computing 10th, 90th percentiles and replacing the outliers

def handle_outliers(data):

    tenth_percentile = np.percentile(data, 10)
    ninetieth_percentile = np.percentile(data, 90)

    b = np.where(data<tenth_percentile, tenth_percentile, data)

    b1 = np.where(b>ninetieth_percentile, ninetieth_percentile, b)
    return b1


for i in ['trestbps', 'chol', 'thalach', 'oldpeak']:
    df[i]=handle_outliers(df[i])

```

```

[15]: # verifying again with boxplot
plt.figure(figsize=(10,6))
df.boxplot(column=['age', 'trestbps', 'chol', 'thalach', 'oldpeak']);

```



```
[16]: # now there are no outliers
```

```
[17]: df.describe().T
```

```
[17]:
```

	count	mean	std	min	25%	50%	75%	max
age	302.0	54.420530	9.047970	29.0	48.00	55.5	61.00	77.0
sex	302.0	0.682119	0.466426	0.0	0.00	1.0	1.00	1.0
cp	302.0	0.963576	1.032044	0.0	0.00	1.0	2.00	3.0
trestbps	302.0	130.523179	13.620063	110.0	120.00	130.0	140.00	152.0
chol	302.0	244.696358	39.237586	188.4	211.00	240.5	274.75	308.9
fbs	302.0	0.149007	0.356686	0.0	0.00	0.0	0.00	1.0
restecg	302.0	0.526490	0.526027	0.0	0.00	1.0	1.00	2.0
thalach	302.0	149.992715	19.608496	116.0	133.25	152.5	166.00	176.8
exang	302.0	0.327815	0.470196	0.0	0.00	0.0	1.00	1.0
oldpeak	302.0	0.966556	0.976726	0.0	0.00	0.8	1.60	2.8
slope	302.0	1.397351	0.616274	0.0	1.00	1.0	2.00	2.0
ca	302.0	0.718543	1.006748	0.0	0.00	0.0	1.00	4.0
thal	302.0	2.314570	0.613026	0.0	2.00	2.0	3.00	3.0
target	302.0	0.543046	0.498970	0.0	0.00	1.0	1.00	1.0

```
[18]: mean_value = df.mean()
median_value = df.median()

std_deviation = df.std()
interquartile_range = df.quantile(0.75) - df.quantile(0.25)

print("Mean:\n", mean_value)
```

Mean:

```
age          54.420530
sex           0.682119
cp            0.963576
trestbps     130.523179
chol         244.696358
fbs           0.149007
restecg       0.526490
thalach      149.992715
exang         0.327815
oldpeak       0.966556
slope         1.397351
ca            0.718543
thal          2.314570
target        0.543046
dtype: float64
```

```
[19]: print("\nMedian:\n", median_value)
```

```
Median:
  age      55.5
  sex       1.0
  cp        1.0
  trestbps 130.0
  chol     240.5
  fbs       0.0
  restecg   1.0
  thalach   152.5
  exang      0.0
  oldpeak   0.8
  slope     1.0
  ca        0.0
  thal      2.0
  target    1.0
dtype: float64
```

```
[20]: print("\nStandard Deviation:\n", std_deviation)
```

```
Standard Deviation:
  age      9.047970
  sex      0.466426
  cp       1.032044
  trestbps 13.620063
  chol     39.237586
  fbs      0.356686
  restecg  0.526027
  thalach  19.608496
  exang     0.470196
  oldpeak  0.976726
  slope    0.616274
  ca       1.006748
  thal     0.613026
  target   0.498970
dtype: float64
```

```
[21]: print("\nInterquartile Range:\n", interquartile_range)
```

```
Interquartile Range:
  age      13.00
  sex       1.00
  cp        2.00
  trestbps 20.00
```

```
chol      63.75
fbs       0.00
restecg   1.00
thalach   32.75
exang     1.00
oldpeak   1.60
slope     1.00
ca        1.00
thal      1.00
target    1.00
dtype: float64
```

3 Exploring categorical features:


Various types of categorical structures are :

```
['sex','cp','fbs','restecg','exang','slope','ca','thal','target']
```

```
[22]: df["sex"].value_counts()
```

```
[22]: 1    206
      0    96
      Name: sex, dtype: int64
```

```
[23]: plt.figure(figsize=(3,3))
      sns.countplot(x = "sex",data = df) # countplot
      plt.show()
```




```
[24]: df["cp"].value_counts()
```



```
[24]: 0    143  
      2    86  
      1    50  
      3    23  
      Name: cp, dtype: int64
```


```
[25]: plt.figure(figsize=(3,3))  
      sns.countplot(x = "cp",data = df)  
      plt.show()
```



```
[26]: df["fbs"].value_counts()
```

```
[26]: 0    257  
      1    45  
      Name: fbs, dtype: int64
```


```
[27]: plt.figure(figsize=(3,3))  
      sns.countplot(x = "fbs",data = df)  
      plt.show()
```



```
[28]: df["restecg"].value_counts()
```

```
[28]: 1    151
      0    147
      2     4
      Name: restecg, dtype: int64
```

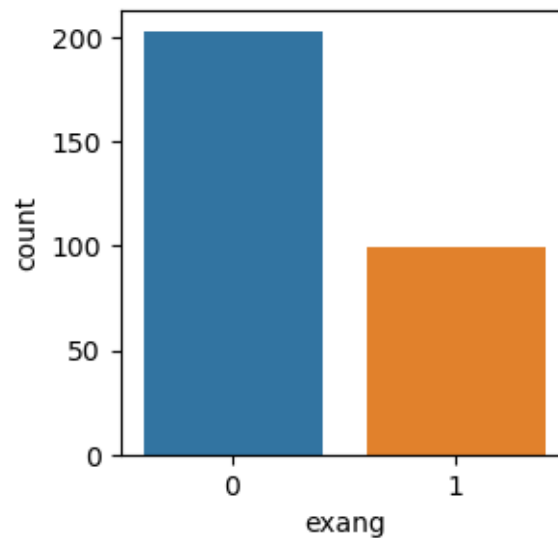
```
[29]: plt.figure(figsize=(3,3))
      sns.countplot(x = "restecg",data = df)
      plt.show()
```



```
[30]: df["exang"].value_counts()
```

```
[30]: 0    203  
      1     99  
      Name: exang, dtype: int64
```


```
[31]: plt.figure(figsize=(3,3))  
      sns.countplot(x = "exang",data = df)  
      plt.show()
```



```
[32]: df["slope"].value_counts()
```

```
[32]: 2    141  
      1    140  
      0     21  
      Name: slope, dtype: int64
```


```
[33]: plt.figure(figsize=(3,3))  
      sns.countplot(x = "slope",data = df)  
      plt.show()
```



```
[34]: df["ca"].value_counts()
```

```
[34]: 0    175  
      1     65  
      2     38  
      3     20  
      4      4  
      Name: ca, dtype: int64
```


```
[35]: plt.figure(figsize=(3,3))  
      sns.countplot(x = "ca",data = df)  
      plt.show()
```



```
[36]: df["thal"].value_counts()
```

```
[36]: 2    165  
      3    117  
      1     18  
      0      2  
      Name: thal, dtype: int64
```


```
[37]: plt.figure(figsize=(3,3))  
      sns.countplot(x = "thal",data = df)  
      plt.show()
```




```
[38]: df["target"].value_counts()
```

```
[38]: 1    164  
      0    138  
      Name: target, dtype: int64
```

```
[39]: plt.figure(figsize=(3,3))  
      sns.countplot(x = "target",data = df)  
      plt.show()
```




```
[40]: # Example: Histogram of Age, separated by CVD status
plt.figure(figsize=(10, 4))
sns.histplot(df, x='age', hue='target', bins=20, kde=True, multiple='stack')
plt.title('Distribution of Age by Cardiovascular Disease (CVD) Status')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```




```
[41]: # Example: Boxplot of Age vs. Target (CVD)
plt.figure(figsize=(10, 4))
sns.boxplot(x='target', y='age', data=df)
```

```
plt.title('Distribution of Age by Cardiovascular Disease (CVD) Status')
plt.xlabel('CVD Status (0: No CVD, 1: CVD)')
plt.ylabel('Age')
plt.show()
```




```
[42]: # Countplot of Sex
plt.figure(figsize=(4, 4))
sns.countplot(x='sex', data=df)
plt.title('Distribution of Patients by Sex')
plt.xlabel('Sex (0: Female, 1: Male)')
plt.ylabel('Count')
plt.show()
```



```
[43]: sex_distribution = df['sex'].value_counts()
labels = ['Male', 'Female']

plt.figure(figsize=(4, 4))
plt.pie(sex_distribution, labels=labels, autopct='%1.1f%%', startangle=90,
        colors=['skyblue', 'lightcoral'])
plt.title('Composition of Patients by Sex')
plt.show()
```


Composition of Patients by Sex




Composition of all patients with respect to the Sex category : From the count plot and histogram we can conclude that There is more number of male patient present compared to female patients.

```
[44]: df_new = df[["trestbps", "target"]].groupby(["trestbps", "target"]).size().  
      ↪reset_index().rename(columns = {0: "number"})  
df_new.head()
```

```
[44]:   trestbps  target  number  
0    110.0      0      15  
1    110.0      1      24  
2    112.0      0       4  
3    112.0      1       5  
4    114.0      0       1
```

```
[45]: # Example: Histogram of trestbps, separated by CVD status  
plt.figure(figsize=(10, 4))  
sns.histplot(df, x='trestbps', hue='target', bins=20, kde=True, multiple='stack')  
plt.show()
```



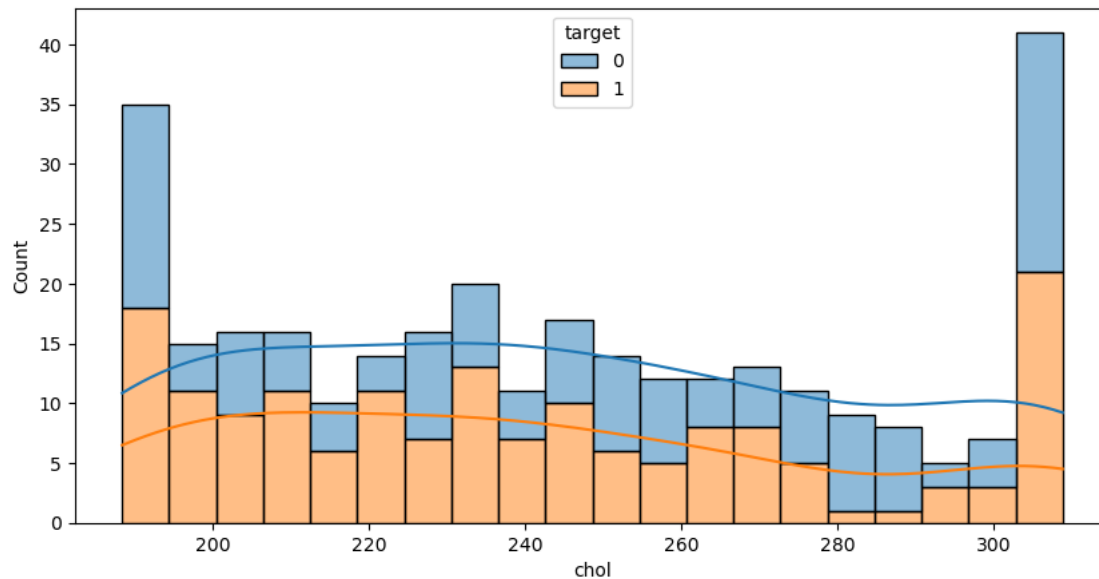
Relationship between resting blood sugar level and target variable : From the histplot we can conclude that a high amount of patient count present in "resting blood sugar level:110", "resting blood sugar level:120", "resting blood sugar level:130", "resting blood sugar level:140" and "resting blood sugar level:150" and hence there is also high amount of share of "occurrence of CVD" is present in the same "resting blood sugar level" also.

```
[46]: df_new_one = df[["chol", "target"]].groupby(["chol", "target"]).size().
      ↪reset_index().rename(columns = {0: "number"})
df_new_one.head(10)
```

```
[46]:
```

	chol	target	number
0	188.4	0	16
1	188.4	1	15
2	192.0	1	2
3	193.0	0	1
4	193.0	1	1
5	195.0	1	1
6	196.0	1	2
7	197.0	0	2
8	197.0	1	4
9	198.0	0	1

```
[47]: # Example: Histogram of chol, separated by CVD status
plt.figure(figsize=(10, 5))
sns.histplot(df, x='chol', hue='target', bins=20, kde=True, multiple='stack')
plt.show()
```




Relationship between cholesterol level and target variable : From the histplot we can conclude that all ranges of cholesterol level have an approximately equal sharing of "occurrence of CVD" though the "occurrence of CVD" is minimum between the range "270 to 305".

```
[48]: df_new_one = df[["slope", "target"]].groupby(["slope", "target"]).size().
      ↪reset_index().rename(columns = {0: "number"})
df_new_one.head()
```


```
[48]:   slope  target  number
0      0        0       12
1      0        1        9
2      1        0       91
3      1        1       49
4      2        0       35
```

```
[49]: plt.figure(figsize=(10, 4))
sns.histplot(df, x='slope', hue='target', bins=20, kde=True, multiple='stack')
plt.show()
```



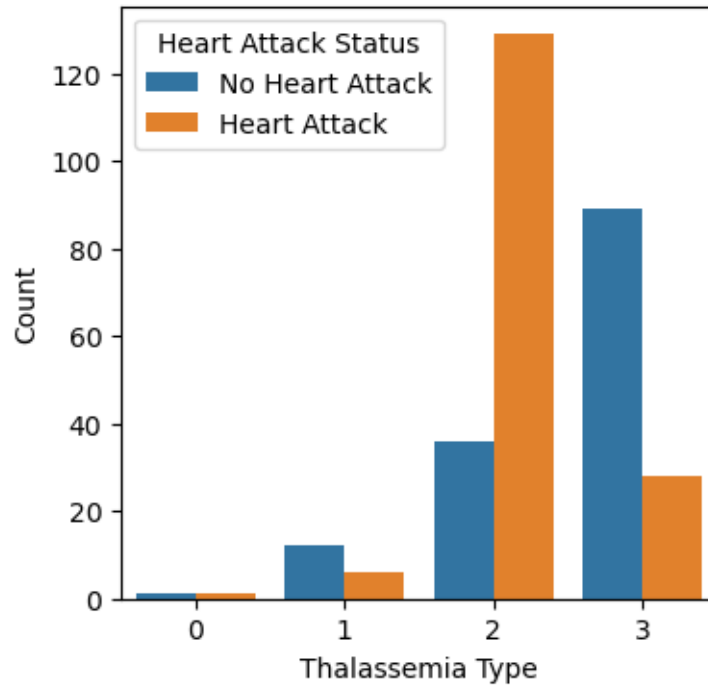
Relationship between peak exercising and occurrence of a heart attack : The variable "slope" which denotes the the peak exercising which contains three classes 0 ,1 ,2 . The class 0 has an approximate share of "occurrence of CVD" . The class 1 has maximum share of "non occurrence of CVD" and the class 2 has a maximum share of "occurrence of CVD".

```
[50]: # Countplot of slope vs. Target (heart attack)
plt.figure(figsize=(4, 4))
sns.countplot(x='slope', hue='target', data=df)
plt.title('Distribution of Slope by Heart Attack Status')
plt.xlabel('Slope of the Peak Exercise ST Segment')
plt.ylabel('Count')
plt.legend(title='Heart Attack Status', labels=['No Heart Attack', 'Heart Attack'])
plt.show()
```



```
[51]: # Countplot of thal vs. Target (heart attack)
plt.figure(figsize=(4, 4))
sns.countplot(x='thal', hue='target', data=df)
plt.title('Distribution of Thalassemia by Heart Attack Status')
plt.xlabel('Thalassemia Type')
plt.ylabel('Count')
plt.legend(title='Heart Attack Status', labels=['No Heart Attack', 'Heart_
Attack'])
plt.show()
```

Distribution of Thalassemia by Heart Attack Status



Relationship between thalassemia type and occurrence of a heart attack : Type 2 and type 3 thalassemia has the majority share of the count . We can conclude from the count plot is that type 2 thalassemia patients have a high number of CVD affected patients with respect to type 3 patients which has a low number of CVD affected patients.

```
[52]: df_tt = df[["target", "thal"]].groupby(["thal"]).value_counts().reset_index().
      ↪ rename(columns = {0: "number"})
      df_tt
```

```
[52]:   thal  target  number
0     0       0        1
1     0       1         1
2     1       0        12
3     1       1         6
4     2       1       129
5     2       0        36
6     3       0        89
7     3       1        28
```

```
[53]: df.head()
```


```
[53]:   age  sex  cp  trestbps   chol  fbs  restecg  thalach  exang  oldpeak  \
0   63   1   3   145.0  233.0   1       0   150.0     0     2.3
1   37   1   2   130.0  250.0   0       1   176.8     0     2.8
2   41   0   1   130.0  204.0   0       0   172.0     0     1.4
```

3	56	1	1	120.0	236.0	0	1	176.8	0	0.8
4	57	0	0	120.0	308.9	0	1	163.0	1	0.6

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1


```
[54]: ## Other factors determining the occurrence of heart attack
```

```
[99]: #for age
plt.figure(figsize=(10, 5))
sns.histplot(df, x='age', hue='target', bins=20, kde=True, multiple='stack')
plt.show()
```




```
[56]: plt.figure(figsize=(10,3))
sns.kdeplot(x='age', data=df, color='red', hue = "target", fill=True, alpha=0.3)
```

```
[56]: <Axes: xlabel='age', ylabel='Density'>
```




Relationship of age and the target variable : patients with the age 41 to 68 have the heighter number of counts and patients with the age 51 to 68 have the maximum number of share of patients with CVD.

```
[57]: # for cp
plt.figure(figsize=(3,3))
sns.countplot(x = "cp",data = df,hue = "target")
plt.show()
```




```
[58]: plt.figure(figsize=(5,3))
sns.kdeplot(x='cp', data=df, color='red',hue = "target", fill=True, alpha=0.3)
```

```
[58]: <Axes: xlabel='cp', ylabel='Density'>
```


Relationship between chest pain type and the target variable : chest pain type 0 has the heighest number of patients and a maximum share of patients without CVD and chest pain type 1 , chest pain type 2, chest pain type 3 has a maximum share of patients with CVD.

```
[59]: ## for fbs
plt.figure(figsize=(15,3))
sns.countplot(x = "fbs",data = df,hue = "target")
plt.show()
```




```
[60]: ## for rest ecg

plt.figure(figsize=(15,3))
sns.countplot(x = "restecg",data = df,hue = "target")
plt.show()
```




```
[61]: plt.figure(figsize=(8,3))
sns.kdeplot(x='restecg', data=df, color='red',hue = "target", fill=True,α
↪alpha=0.3)
```

```
[61]: <Axes: xlabel='restecg', ylabel='Density'>
```

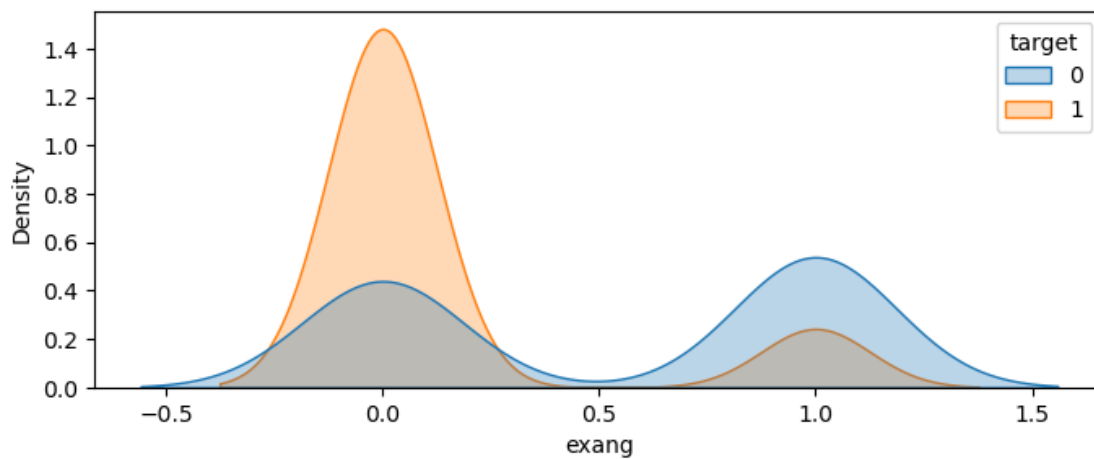


```
[101]: # for exang
plt.figure(figsize=(10, 4))
sns.histplot(df, x='exang',hue='target', bins=20, kde=True, multiple='stack')
plt.show()
```




```
[65]: plt.figure(figsize=(8,3))
sns.kdeplot(x='exang', data=df, color='red',hue = "target", fill=True, alpha=0.
↪3)
```

```
[65]: <Axes: xlabel='exang', ylabel='Density'>
```




```
[102]: # for oldpeak
plt.figure(figsize=(10, 4))
sns.histplot(df, x='oldpeak',hue='target', bins=20, kde=True, multiple='stack')
plt.show()
```



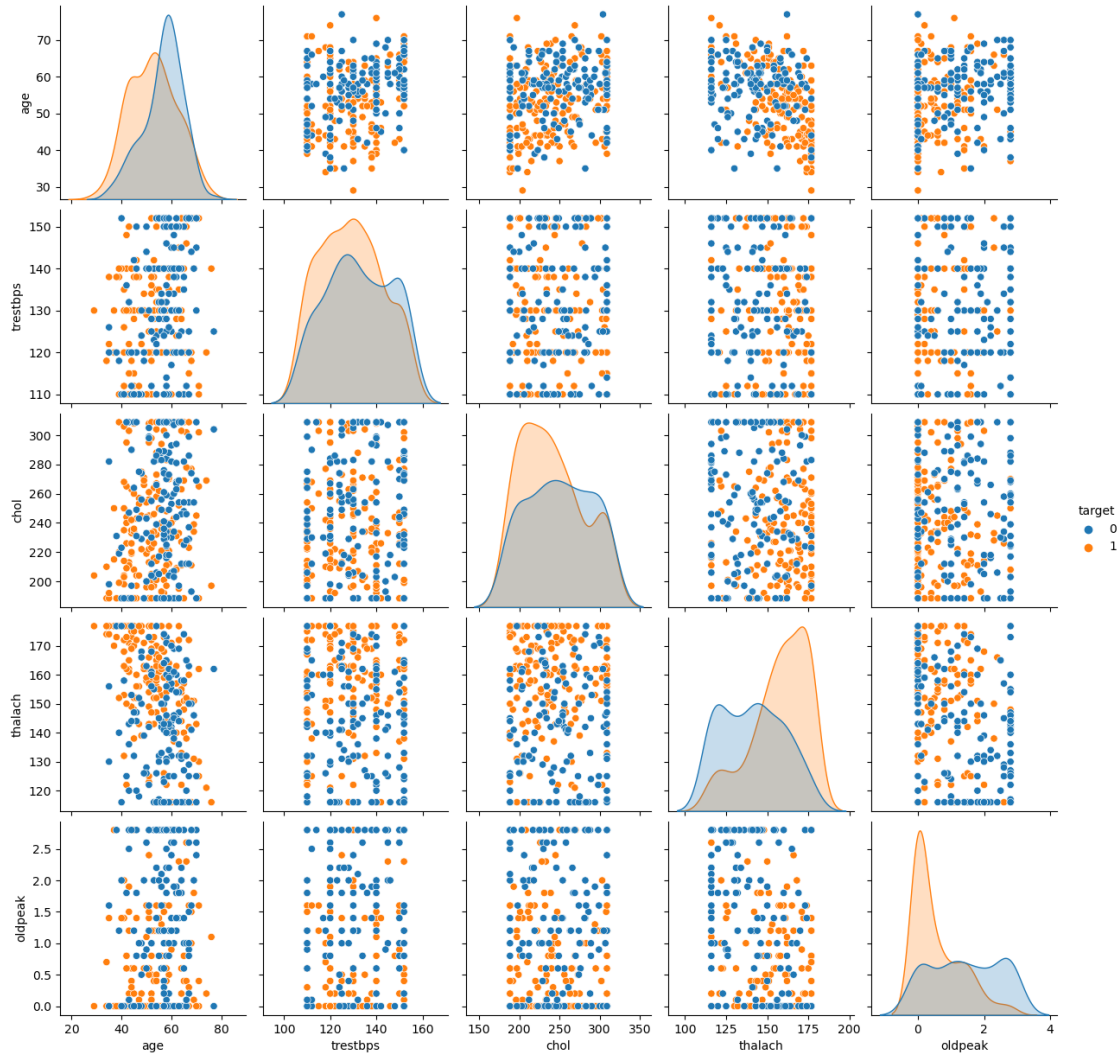
```
[67]: plt.figure(figsize=(8,3))
sns.kdeplot(x='oldpeak', data=df,hue = "target", fill=True, alpha=0.3)
```

```
[67]: <Axes: xlabel='oldpeak', ylabel='Density'>
```



```
[105]: ## pairplot
# pairplot for numeric columns
sns.
    pairplot(df[['age', 'trestbps', 'chol', 'thalach', 'oldpeak', 'target']], hue='target')
```

```
[105]: <seaborn.axisgrid.PairGrid at 0x1a604a72bc0>
```




3.1 Correlation analysis

```
[129]: # Compute the correlation matrix
correlation_matrix = df.corr()

# Create a heatmap using Seaborn
plt.figure(figsize=(12, 6)) # Set the figure size
# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(correlation_matrix, annot=True)

plt.title("Correlation Matrix Heatmap")
plt.show()
```



3.2 Feature Scaling

```
[71]: ## Feature Scaling
```

```
[108]: from sklearn.preprocessing import StandardScaler
# -infinity to infinity
# z = (X - mean)/std deviation

scale = StandardScaler()
df[['age', 'trestbps', 'chol', 'thalach']] = scale.
    ↪ fit_transform(df[['age', 'trestbps', 'chol', 'thalach']])
```

3.3 Preparing The data

```
[109]: x = df.drop(["target"], axis = 1)
y_init = df["target"]
```

```
[110]: y = y_init.to_numpy()
```

```
[111]: from sklearn.model_selection import train_test_split
```

```
[112]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.
    ↪ 2, random_state = 1) # 80 : 20
```

```
[113]: x_train.shape
```

```
[113]: (241, 13)
```

```
[114]: y_train.shape
```

```
[114]: (241,)
```

4 2. Model building

4.1 Apply logistic regression and evaluating performance

```
[115]: from sklearn.linear_model import LogisticRegression
```

```
[116]: LR_model = LogisticRegression()
```

```
[117]: LR_model.fit(x_train,y_train)
```

```
[117]: LogisticRegression()
```

```
[118]: y_pred = LR_model.predict(x_test)
```

```
[119]: from sklearn.metrics import accuracy_score, classification_report ,  
      ↪confusion_matrix  
accuracy_logreg = accuracy_score(y_test, y_pred) * 100  
print(f'Accuracy of Logistic Regression: {accuracy_logreg:.4f}')
```

`print("\n")`

Additional evaluation metrics

```
print(classification_report(y_test, y_pred))
```

Accuracy of Logistic Regression: 80.3279


	precision	recall	f1-score	support
0	0.84	0.72	0.78	29
1	0.78	0.88	0.82	32
accuracy			0.80	61
macro avg	0.81	0.80	0.80	61
weighted avg	0.81	0.80	0.80	61

```
[121]: # Compute the confusion matrix  
cm = confusion_matrix(y_test, y_pred)  
# Visualize the confusion matrix using Seaborn  
sns.set(font_scale=1.2) # Adjust font size for better readability
```

```

sns.heatmap(cm, annot=True, fmt="d", cbar=False,
            xticklabels=["Class 0", "Class 1"],
            yticklabels=["Class 0", "Class 1"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.show()

```



```

[122]: import statsmodels.api as sm

# Add a constant term to the features matrix
X_train_with_const = sm.add_constant(x_train)

# Fit logistic regression using statsmodels
logit_model = sm.Logit(y_train, X_train_with_const)
result = logit_model.fit()

# Display summary including coefficients and p-values
print(result.summary())

```


Optimization terminated successfully.
 Current function value: 0.337416
 Iterations 7

Logit Regression Results

Dep. Variable:	y	No. Observations:	241
Model:	Logit	Df Residuals:	227
Method:	MLE	Df Model:	13
Date:	Wed, 20 Dec 2023	Pseudo R-squ.:	0.5100
Time:	11:49:51	Log-Likelihood:	-81.317
converged:	True	LL-Null:	-165.95
Covariance Type:	nonrobust	LLR p-value:	2.605e-29

	coef	std err	z	P> z	[0.025	0.975]
const	3.3485	1.111	3.013	0.003	1.170	5.526
age	0.0351	0.234	0.150	0.881	-0.424	0.494
sex	-1.9547	0.552	-3.543	0.000	-3.036	-0.873
cp	0.9485	0.208	4.558	0.000	0.541	1.356
trestbps	-0.2696	0.207	-1.302	0.193	-0.676	0.136
chol	-0.4675	0.243	-1.928	0.054	-0.943	0.008
fbs	0.2282	0.621	0.367	0.713	-0.989	1.445
restecg	0.1490	0.399	0.374	0.709	-0.633	0.931
thalach	0.5510	0.262	2.104	0.035	0.038	1.064
exang	-0.8479	0.478	-1.774	0.076	-1.785	0.089
oldpeak	-0.6882	0.270	-2.553	0.011	-1.217	-0.160
slope	0.4790	0.406	1.179	0.238	-0.317	1.275
ca	-0.9820	0.248	-3.956	0.000	-1.468	-0.495
thal	-0.8409	0.328	-2.562	0.010	-1.484	-0.198

4.2 Applying Random Forest and evaluating performance

```
[130]: from sklearn.ensemble import RandomForestClassifier
my_rf_classifier = RandomForestClassifier()
```

```
[131]: my_rf_classifier.fit(x_train, y_train)
```

```
[131]: RandomForestClassifier()
```

```
[132]: my_predictions = my_rf_classifier.predict(x_test)
```


```
[135]: print(accuracy_score(y_test, my_predictions)*100)
```

80.32786885245902

```
[136]: print(classification_report(y_test, my_predictions))
```

	precision	recall	f1-score	support
0	0.84	0.72	0.78	29
1	0.78	0.88	0.82	32
accuracy			0.80	61
macro avg	0.81	0.80	0.80	61
weighted avg	0.81	0.80	0.80	61

```
[137]: # Compute the confusion matrix
cm = confusion_matrix(y_test, my_predictions)
# Visualize the confusion matrix using Seaborn
sns.set(font_scale=1.2) # Adjust font size for better readability
sns.heatmap(cm, annot=True, fmt="d", cbar=False,
            xticklabels=["Class 0", "Class 1"],
            yticklabels=["Class 0", "Class 1"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.show()
```



4.2.1 Hyperparameter tuning:

```
[175]: my_rf_classifier1 = RandomForestClassifier(n_estimators=50,criterion='entropy')
my_rf_classifier1.fit(x_train, y_train)
my_predictions1 = my_rf_classifier1.predict(x_test)
print(accuracy_score(y_test, my_predictions1))
```

0.819672131147541

```
[176]: print(classification_report(y_test, my_predictions1))
```

	precision	recall	f1-score	support
0	0.85	0.76	0.80	29
1	0.80	0.88	0.84	32
accuracy			0.82	61
macro avg	0.82	0.82	0.82	61
weighted avg	0.82	0.82	0.82	61

```
[ ]:
```